


Same graph, two control flows

GRAPH-MCP (tool)



ranking happens INSIDE the model's context → under-ranks

GRAPH-NATIVE (harness)



ranking happens IN CODE before the agent thinks → keeps the right few

↑ the ONLY difference between the arms ↑

The ranker (the algorithm)

score(file) =

reference density <i>how many times it uses the symbol</i>	+refs
direct-caller bonus <i>1-hop caller (additive, not a tier)</i>	+8
name-match <i>basename⇒anchor → subtype/impl</i>	+10
package locality <i>same package tiebreak</i>	+4
TEST-FILE penalty <i>Test*/!Test*/ /test/ → buried</i>	-10⁵

sort ↓ score, cap top-20. Held-out F_1 : raw 0.17 → ranked 0.52